

ARTICLES CATEGORIZER

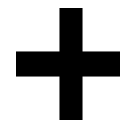
DATA MINING AND MACHINE LEARNING
MIRKO DI LUCIA



UNIVERSITÀ DI PISA

INTRODUCTION

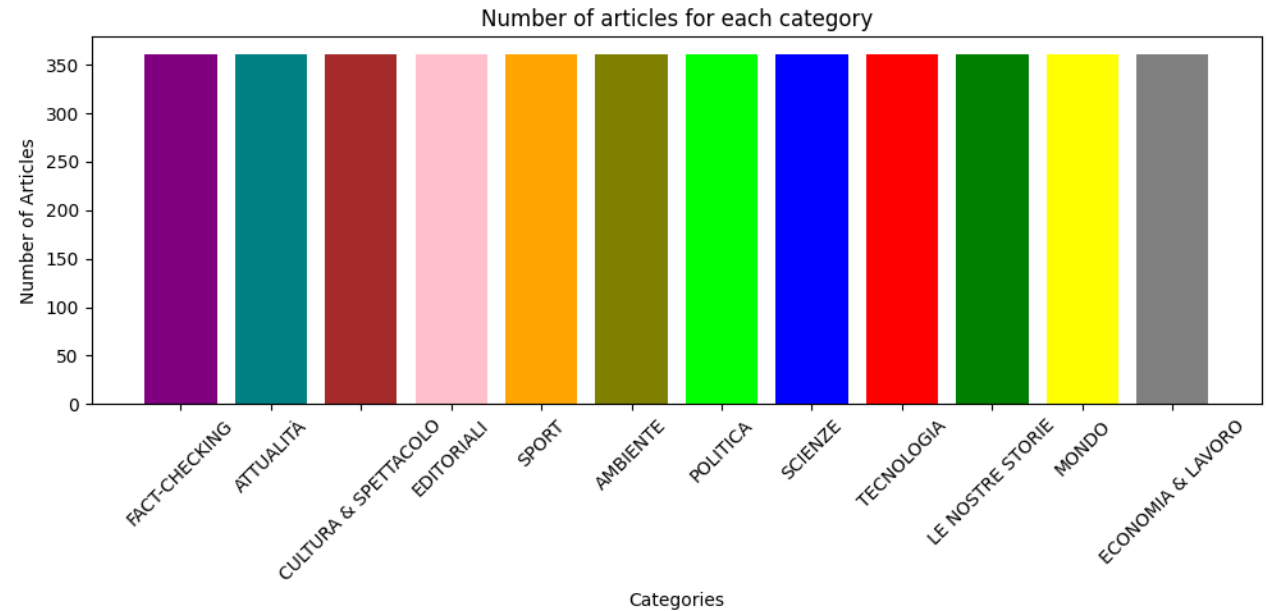
Find the most reliable data mining algorithm, and structure to categorize a subgroup of articles from a famous online newspaper website based on category



DATASET DESCRIPTION

Dataset was scraped from the website

- Dataset is composed from 4332 articles
- Each article had a single category



DATASET DESCRIPTION

Top commons words in each category

- FACT-CHECKING: (stat, 1408), (articol, 919), (facebook, 907), (vide, 879), (qui, 848), (legg, 823), (post, 768), (pubblic, 712), (vaccin, 672), (fals, 618)
- ATTUALITÀ: (stat, 950), (cas, 396), (prim, 386), (poi, 355), (pi, 320), (fatt, 286), (trov, 284), (anni, 266), (person, 265), (due, 259)
- CULTURA & SPETTACOLO: (stat, 592), (pi, 527), (cos, 447), (prim, 375), (anni, 338), (rai, 323), (due, 321), (pubblic, 320), (poi, 319), (perc, 317)
- EDITORIALI: (pi, 518), (prim, 508), (salvin, 495), (part, 489), (cos, 482), (fatt, 460), (stat, 451), (polit, 433), (govern, 421), (ital, 337)
- SPORT: (stat, 603), (atlet, 490), (parig, 414), (final, 394), (dop, 384), (part, 365), (prim, 349), (olimpiad, 332), (pi, 331), (gar, 297)
- AMBIENTE: (pi, 1181), (stat, 805), (climat, 533), (paes, 446), (camb, 431), (ital, 409), (nuov, 407), (esser, 376), (europe, 374), (second, 368)
- POLITICA: (stat, 770), (part, 696), (ital, 599), (pi, 591), (polit, 463), (govern, 422), (prim, 391), (president, 380), (melon, 361), (sol, 325)
- SCIENZE: (pi, 1005), (stat, 940), (ricerc, 615), (stud, 593), (vaccin, 588), (prim, 581), (esser, 504), (cas, 489), (nuov, 464), (sol, 448)
- TECNOLOGIA: (pi, 628), (stat, 598), (utent, 371), (esser, 342), (dat, 341), (nuov, 340), (piattaform, 327), (prim, 320), (propr, 314), (artificial, 304)
- LE NOSTRE STORIE: (pi, 1766), (stat, 1365), (lavor, 1181), (cos, 1043), (prim, 967), (anni, 860), (cas, 813), (sol, 773), (perc, 772), (far, 749)
- MONDO: (stat, 1124), (pi, 461), (israel, 448), (prim, 439), (russ, 403), (part, 393), (second, 381), (president, 341), (harris, 337), (dop, 331)
- ECONOMIA & LAVORO: (pi, 960), (lavor, 754), (stat, 639), (eur, 618), (ital, 583), (anni, 462), (europe, 454), (part, 445), (prim, 432), (nuov, 418)

THE SCRAPER PIPELINE

Fetch
categories url

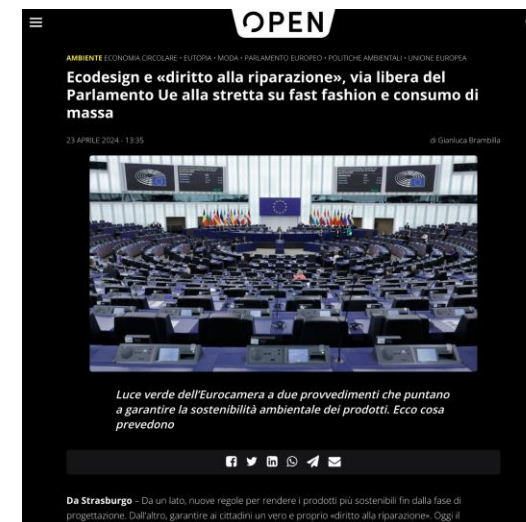
1

Fetch all articles url
in each category

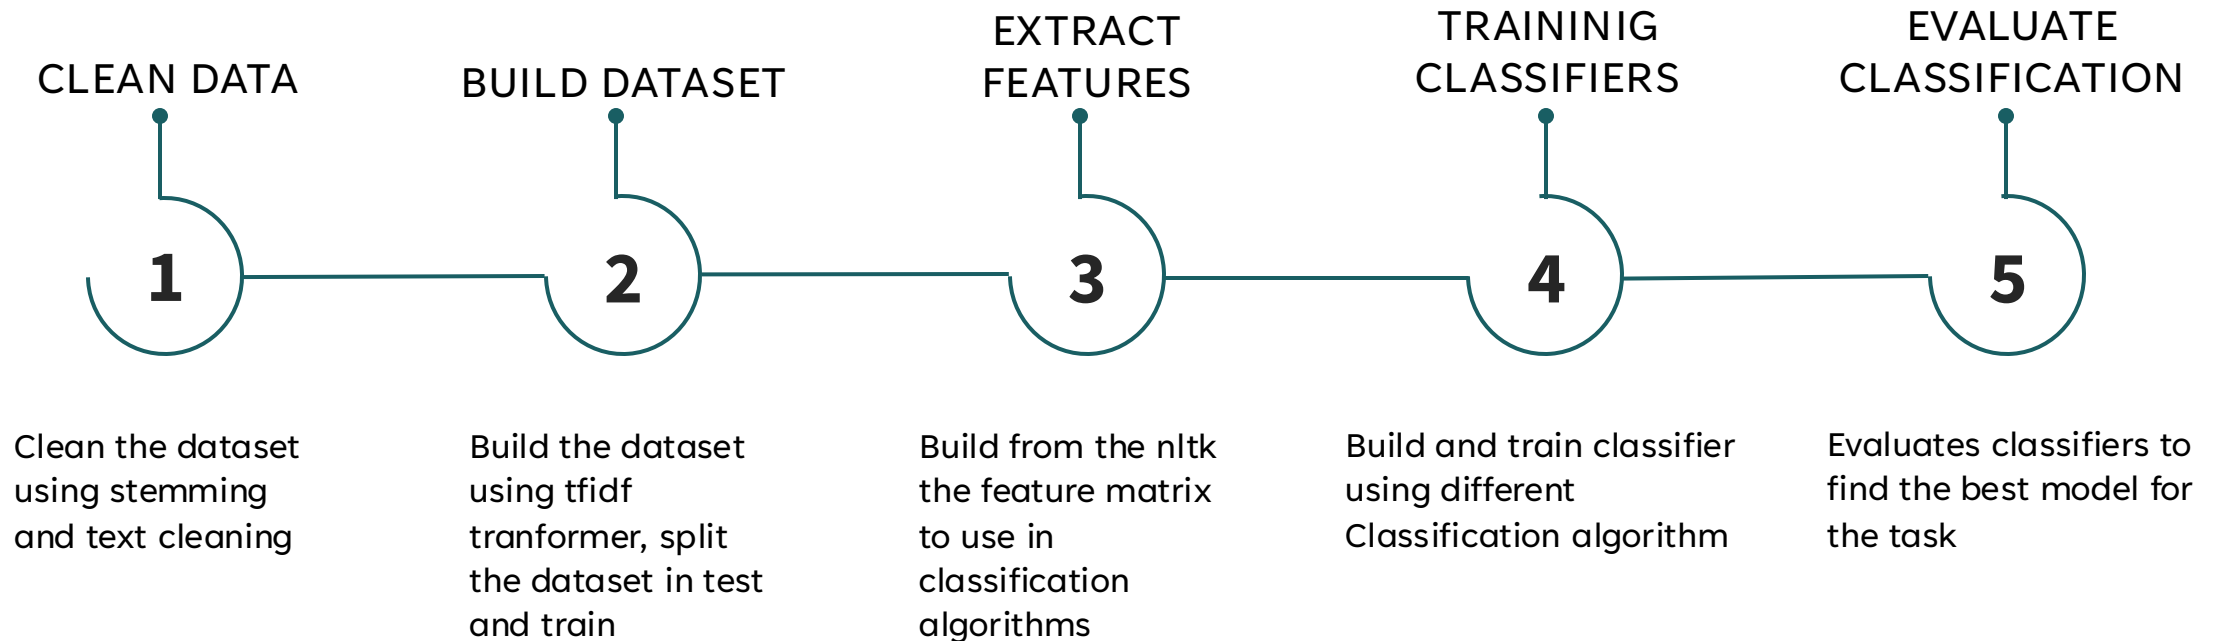
2

Scrape all
articles

3

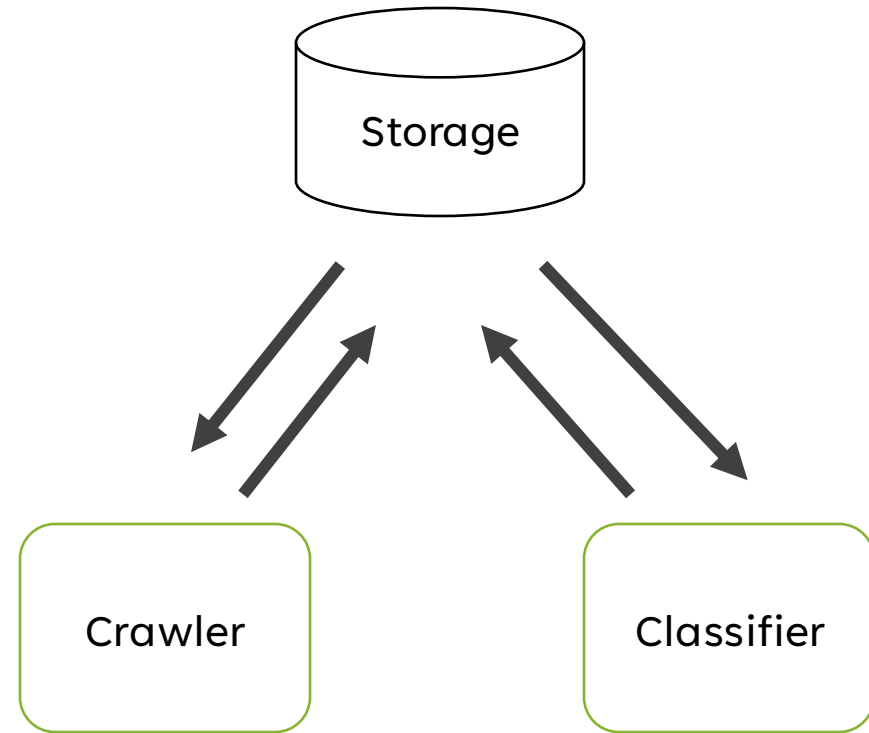


TRAINING PIPELINE



APPLICATION STRUCTURE

- Application is a series of python scripts divided in 2 main folders
 - Crawler
 - Classifier



+



CLASSIFIERS

- Decision tree
- KNeighbors
- Multinomial Naive Bayes
- Random Forest

RESULTS - Decision Tree

Accuracy on tests set:

0.7707317073170732

Metrics per class on tests set:

	precision	recall	f1-score	support
<u>AMBIENTE</u>	0.74	0.76	0.75	70
<u>ATTUALITÀ</u>	0.96	0.96	0.96	68
<u>CULTURA & SPETTACOLO</u>	0.96	0.97	0.97	72
<u>ECONOMIA & LAVORO</u>	0.64	0.62	0.63	55
<u>MONDO</u>	0.61	0.63	0.62	84
<u>POLITICA</u>	0.83	0.89	0.86	64
<u>SCIENZE</u>	0.72	0.72	0.72	64
<u>SPORT</u>	0.83	0.75	0.79	67
<u>TECNOLOGIA</u>	0.67	0.65	0.66	71
accuracy			0.77	615
macro avg	0.77	0.77	0.77	615
weighted avg	0.77	0.77	0.77	615

Confusion matrix:

	amb	att	c&s	eco	mon	pol	sci	spo	tec
amb	[53	0	0	3	2	1	5	2	4]
att	[1	65	0	0	0	1	0	0	1]
c&s	[0	0	70	0	0	0	0	0	2]
eco	[3	0	1	34	7	1	3	0	6]
mon	[7	1	0	2	53	7	4	5	5]
pol	[0	1	0	2	2	57	1	1	0]
sci	[6	1	0	1	5	0	46	2	3]
spo	[1	0	0	2	8	2	2	50	2]
tec	[1	0	2	9	10	0	3	0	46]

RESULTS - Random Forest

Accuracy on tests set:

0.9056910569105691

Metrics per class on tests set:

	precision	recall	f1-score	support
<u>AMBIENTE</u>	0.86	0.90	0.88	70
<u>ATTUALITÀ</u>	0.92	0.96	0.94	68
<u>CULTURA & SPETTACOLO</u>	0.99	1.00	0.99	72
<u>ECONOMIA & LAVORO</u>	0.79	0.91	0.85	55
<u>MONDO</u>	0.94	0.79	0.86	84
<u>POLITICA</u>	0.91	0.92	0.91	64
<u>SCIENZE</u>	0.88	0.88	0.88	64
<u>SPORT</u>	0.94	0.97	0.96	67
<u>TECNOLOGIA</u>	0.91	0.86	0.88	71
accuracy			0.91	615
macro avg	0.90	0.91	0.90	615
weighted avg	0.91	0.91	0.91	615

Confusion matrix:

	amb	att	c&s	eco	mon	pol	sci	spo	tec
amb	[63	0	0	2	1	0	3	0	1]
att	[0	65	0	0	1	1	1	0	0]
c&s	[0	0	72	0	0	0	0	0	0]
eco	[2	2	0	50	0	1	0	0	0]
mon	[2	4	0	1	66	4	1	3	3]
pol	[0	0	0	5	0	59	0	0	0]
sci	[6	0	0	0	0	0	56	1	1]
spo	[0	0	1	0	0	0	0	65	1]
tec	[0	0	0	5	2	0	3	0	61]

RESULTS – Naive Bayes

Accuracy on tests set:

0.8991869918699187

Metrics per class on tests set:

	precision	recall	f1-score	support
<u>AMBIENTE</u>	0.86	0.91	0.89	70
<u>ATTUALITÀ</u>	0.90	0.91	0.91	68
<u>CULTURA & SPETTACOLO</u>	0.95	1.00	0.97	72
<u>ECONOMIA & LAVORO</u>	0.77	0.91	0.83	55
<u>MONDO</u>	0.97	0.76	0.85	84
<u>POLITICA</u>	0.83	0.89	0.86	64
<u>SCIENZE</u>	0.91	0.91	0.91	64
<u>SPORT</u>	0.94	0.96	0.95	67
<u>TECNOLOGIA</u>	0.97	0.87	0.92	71
accuracy			0.90	615
macro avg	0.90	0.90	0.90	615
weighted avg	0.90	0.90	0.90	615

Confusion matrix:

	amb	att	c&s	eco	mon	pol	sci	spo	tec
amb	[64	0	1	1	0	0	3	0	1]
att	[0	62	1	2	1	1	1	0	0]
c&s	[0	0	72	0	0	0	0	0	0]
eco	[2	0	0	50	0	3	0	0	0]
mon	[2	7	0	2	64	6	0	3	0]
pol	[1	0	0	6	0	57	0	0	0]
sci	[5	0	0	0	0	0	58	0	1]
spo	[0	0	2	0	1	0	0	64	0]
tec	[0	0	0	4	0	2	2	1	62]

RESULTS - kNN

Accuracy on tests set:

0.8666666666666667

Metrics per class on tests set:

	precision	recall	f1-score	support
<u>AMBIENTE</u>	0.79	0.87	0.83	70
<u>ATTUALITÀ</u>	0.86	0.84	0.85	68
<u>CULTURA & SPETTACOLO</u>	0.94	1.00	0.97	72
<u>ECONOMIA & LAVORO</u>	0.76	0.82	0.79	55
<u>MONDO</u>	0.87	0.79	0.82	84
<u>POLITICA</u>	0.84	0.91	0.87	64
<u>SCIENZE</u>	0.86	0.84	0.85	64
<u>SPORT</u>	0.93	0.93	0.93	67
<u>TECNOLOGIA</u>	0.95	0.82	0.88	71
accuracy			0.87	615
macro avg	0.87	0.87	0.87	615
weighted avg	0.87	0.87	0.87	615

Confusion matrix:

	amb	att	c&s	eco	mon	pol	sci	spo	tec
amb	[61	1	0	3	0	1	4	0	0]
att	[2	57	4	1	2	1	1	0	0]
c&s	[0	0	72	0	0	0	0	0	0]
eco	[4	0	0	45	0	4	1	0	1]
mon	[3	5	1	1	66	3	2	3	0]
pol	[1	2	0	1	2	58	0	0	0]
sci	[4	1	0	2	1	0	54	1	1]
spo	[2	0	0	1	1	0	0	62	1]
tec	[0	0	0	5	4	2	1	1	58]

Cross-Validation

Decision Tree Results:

- Decision Tree Cross-Validation Scores: [0.79545455, 0.78030303, 0.77272727, 0.83333333, 0.76425856, 0.74144487, 0.75285171, 0.82509506, 0.79087452, 0.78326996]
- Decision Tree Cross-Validation Mean Accuracy: 0.7839612858624265

Random Forest Results:

- Random Forest Cross-Validation Scores: [0.90151515, 0.92045455, 0.92045455, 0.92045455, 0.91254753, 0.87452471, 0.88973384, 0.92775665, 0.88212928, 0.90874525]
- Random Forest Cross-Validation Mean Accuracy: 0.9058316050236203

KNN Results:

- Cross-Validation Scores: [0.85606061, 0.82575758, 0.86742424, 0.84848485, 0.87452471, 0.79467681, 0.84410646, 0.84790875, 0.79847909, 0.8365019]
- kNN Cross-Validation Mean Accuracy: 0.8393924991358451

Naïve Bayes

- Naive Bayes Cross-Validation Scores: [0.90151515, 0.89772727, 0.90151515, 0.89015152, 0.91634981, 0.86311787, 0.87072243, 0.92395437, 0.88212928, 0.8973384]
- Naive Bayes Cross-Validation Mean Accuracy: 0.8944521258209471